

Sampling distributions,
standard errors, and confidence intervals

Overview

Review of bias and sampling distributions

Exploring sampling distributions in R and the Standard Error


Point estimates and confidence intervals

Review: sampling and sampling distributions

Review: sampling




1	orange
2	red
3	green
4	white
5	white
6	white
7	white
8	white
9	red



Q: What symbol do we use to denote the sample size?

A: ***n***

Bias and the Gettysburg address word length distribution





Bias and the Gettysburg address word length distribution

Bias is when the average statistic values does not equal the population parameter

Here:

$$E[\bar{x}] \neq \mu$$




Statistical bias

μ



\bar{x}





As part of a strategic-planning process, in spring 2013 Hampshire College launched a survey of alums.

Via email, the College **invited 8,160 alums to fill out an online questionnaire** administered by the campus's offices.

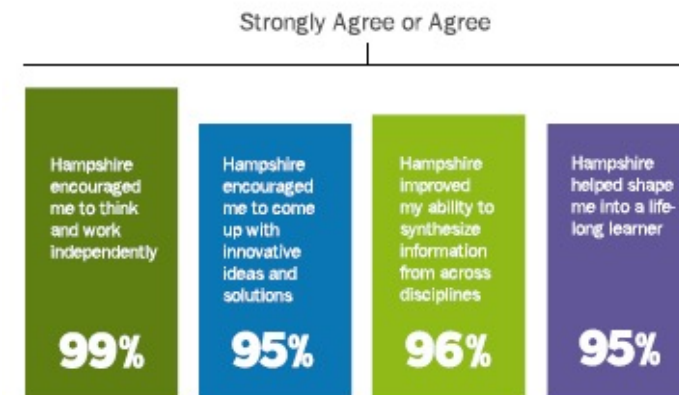
A total of 1,920 surveys were completed, yielding a response rate of 24%.

Alumni Survey Results

As part of a strategic-planning process, in spring 2013 Hampshire College launched a survey of alums. Via email, the College invited 8,160 alums to fill out an online questionnaire administered by the campus's Alumni and Family Relations and Institutional Research offices. A total of 1,920 surveys were completed, yielding a response rate of 24%.

Note: The percentages in the data (below) are based on the number of responses received for each question.

To what extent do you agree with the following statements?



Please rate your student experience at Hampshire.



65% of our alumni earn advanced degrees within ten years of graduating.

1 in 7 alumni holds a Ph.D. or other terminal degree.

Hampshire ranks in the **top 1%** of colleges nationwide in the % of grads that go on to earn doctorates.

26% of our graduates have started their own business or organization.

“

Hampshire does a great job fostering the ability to ask good questions and to look at ideas with a critical lens.


Hampshire has encouraged me to be more engaged, socially aware and more of a critical thinker than my peers.

I feel more able to adapt to a range of environments because Hampshire taught me skills and ideas rather than just knowledge.

”

Bias or No Bias?

$$E[\hat{p}_{\text{replied}}] \neq \pi_{\text{all}}$$



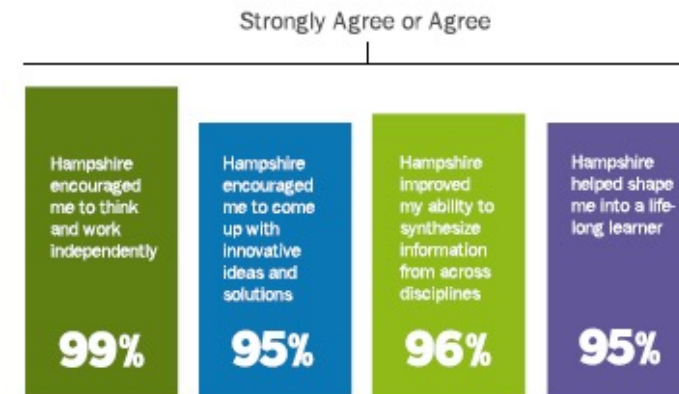
Sad Plato says:
"There's no Truth in advertising"

Alumni Survey Results

As part of a strategic-planning process, in spring 2013 Hampshire College launched a survey of alums. Via email, the College invited 8,160 alums to fill out an online questionnaire administered by the campus's Alumni and Family Relations and Institutional Research offices. A total of 1,920 surveys were completed, yielding a response rate of 24%.

Note: The percentages in the data (below) are based on the number of responses received for each question.

To what extent do you agree with the following statements?



Please rate your student experience at Hampshire.



65% of our alumni earn advanced degrees within ten years of graduating.

1 in 7 alumni holds a Ph.D. or other terminal degree.

Hampshire ranks in the **top 1%** of colleges nationwide in the % of grads that go on to earn doctorates.

26% of our graduates have started their own business or organization.

“

Hampshire does a great job fostering the ability to ask good questions and to look at ideas with a critical lens.

Hampshire has encouraged me to be more engaged, socially aware and more of a critical thinker than my peers.

I feel more able to adapt to a range of environments because Hampshire taught me skills and ideas rather than just knowledge.

”

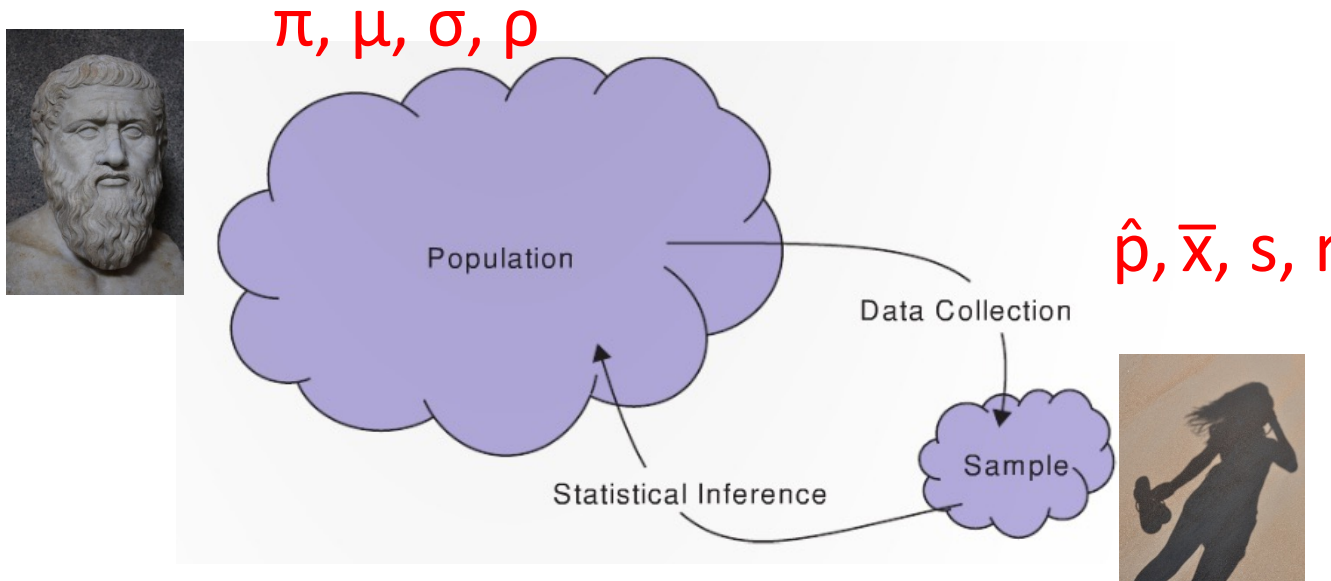
Q: How can we prevent sampling bias?

A: To prevent bias, use a **simple random sample**

- where each member in the population is equally likely to be in the sample

This allows for generalizations to the population!

Soup analogy!



Q: How do we select a random sample?

Mechanically:

- Flip coins

- Pull balls from well mixed bins


- Deal out shuffled cards, etc.

Use a computer program

Q: What computer program can we use?





Q: If you have questions about statistical bias where should you ask them?



A: Ed discussions or come to class!

From now on we are going to assume no bias!

Happy Plato and Lincoln




Our statistic values, on average, reflect the parameters



Exploring sampling distributions in R

For our distribution of Gettysburg word lengths...

The mean length of 10 words (\bar{x}) from different student's samples:





If we had infinite students taking samples yielding infinite statistics, we would have a **distribution** of statistics

Sampling distribution


A **sampling distribution** is the distribution of sample statistics computed for different samples of the same size (n) from the same population

A sampling distribution shows us how the sample statistic varies from sample to sample

Gettysburg address word length sampling distribution



Sampling distribution!



[Gettysburg sampling distribution app](#)

Let's create a sampling distribution in R!

Let's create a sampling distribution in R

Load the SDS100 library to make all SDS100 functions available

```
> library(SDS100)
```

Get the Gettysburg population data

```
> download_data("gettysburg.Rda")
```

```
> load("gettysburg.Rda")
```

```
> word_lengths <- gettysburg$num_letters
```

Let's create a sampling distribution in R

We can use the `sample(data_vec, n)` to get a sample of length `n`:

```
> curr_sample <- sample(word_lengths, 10)
```

Q: How can we get \bar{x} from this sample in R?

```
> mean(curr_sample)
```

Q: How could we get a full sampling distribution?

- A: Repeat this many times to get an approximation of the sampling distribution
- If we store the \bar{x} 's in a vector, we can then plot the sampling distribution as a histogram

The do_it() function

```
do_it(100) * {
```


```
    2 + 3
```

```
}
```

Let's create a sampling distribution in R

```
sampling_dist <- do_it(10000) * {  
  
    curr_sample <- sample(word_lengths, 10)  
    mean(curr_sample)  
  
}  
  
hist(sampling_dist)
```


Sampling distribution in R



`mean(sampling_dist)`

`mean(word_lengths)` # these are the same so no bias

Changing the sample size n

What happens to the sampling distribution as we change n ?


- Experiment for $n = 1, 5, 10, 20$

```
sampling_dist <- do_it(10000) * {  
    curr_sample <- sample(word_lengths, 20)  
    mean(curr_sample)  
}
```


```
hist(sample_means, nclass = 100)
```

[Gettysburg sampling distribution app](#)


Sampling distribution ($n = 1$)




Sampling distribution ($n = 5$)



Sampling distribution ($n = 10$)



Sampling distribution ($n = 20$)



x-axis range 9 vs. 6


As the sample size n increases

1. The sampling distribution becomes more like a normal distribution
2. The sampling distribution points (\bar{x} 's) become more concentrated around the mean $E[\bar{x}] = \mu$


The standard error

The **standard error** of a statistic, denoted SE, is the standard deviation of the sample statistic

- i.e., SE is the standard deviation of the *sampling distribution*



What does the size of a standard error tell us?



Q: If we have a large SE, would we believe a given statistic is a good estimate for the parameter?

- E.g., would we believe a particular \bar{x} is a good estimate for μ ?

A: A large SE means our statistic (point estimate) could be far from the parameter

- E.g., \bar{x} could be far from μ

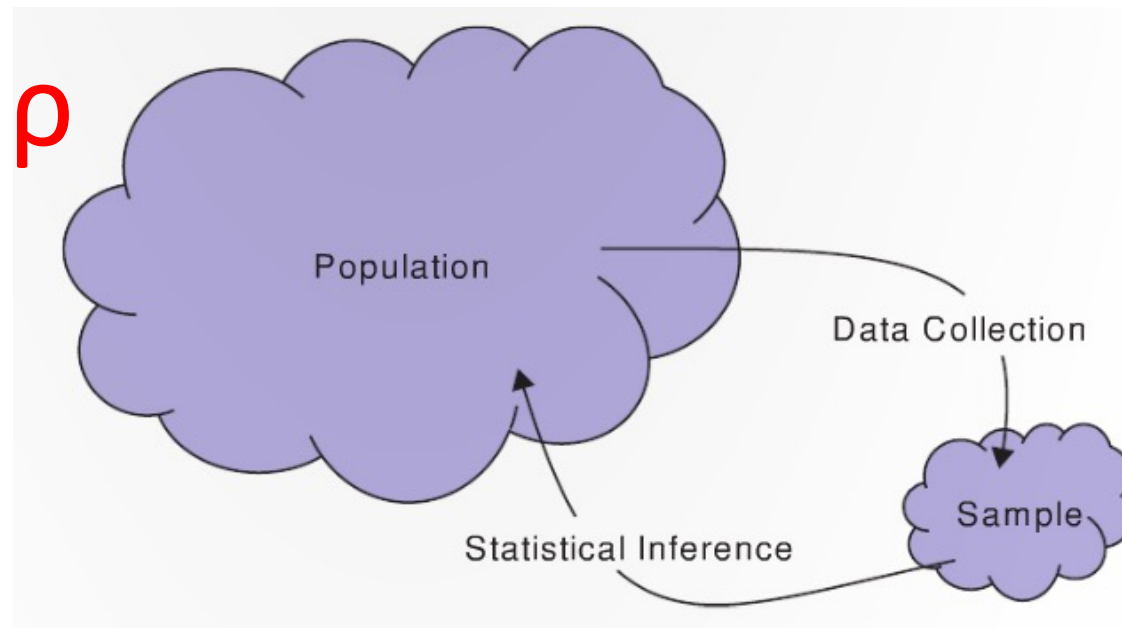

Point estimates and confidence intervals

Back to the big picture: Inference


Statistical inference is...?

the process of drawing conclusions about the entire population based on information in a sample

π, μ, σ, ρ



\hat{p}, \bar{x}, s, r



Point Estimate

We use a statistic from a sample as a **point estimate** for a population parameter

- \bar{x} is a point estimate for...? μ

54% of American approve of Biden's job performance according to a recent Gallup poll

Q: What are π and \hat{p} here?

Q: Is \hat{p} a good estimate for π in this case?

A: We can't tell from the information given

Interval estimate based on a margin of error

An **interval estimate** give a range of plausible values for a population parameter.

One common form of an interval estimate is:

Point estimate \pm margin of error

Where the **margin of error** is a number that reflects the precision of the sample statistic as a point estimate for this parameter

Example: Gallup poll

54% of American approve of Biden's job performance, plus or minus 3%

How do we interpret this?

Says that the population parameter (π) lies somewhere between 51% to 57%

i.e., if they sampled all voters the true population proportion (π) would be likely be in this range

Confidence Intervals

A **confidence interval** is an interval computed by a method that will contain the *parameter* a specified percent of times

- i.e., if the estimation were repeated many times, the interval will have the parameter x% of the time

The **confidence level** is the percent of all intervals that contain the parameter

Think ring toss...

Parameter exists in the ideal world

We toss intervals at it


95% of those intervals capture the parameter




Confidence Intervals

For a **confidence level** of 95%...

95% of the **confidence intervals** will have the parameter in them



Wits and Wagers: 90% confidence interval estimator



I will ask 10 questions that have numeric answers

Please come up with a range of values that contains the true value in it for 9 out of the 10 questions

- i.e., be a 90% confidence interval estimator

Enter your range of estimates for each question below as two numbers separated by a comma

- E.g., 10.2, 50.7

Wits and Wagers...

Question 1: What percent of the world's surface is water?

Question 2: How many floors does the leaning tower of Pisa have?

Question 3: What year was the parking meter invented?

Wits and Wagers...

Question 4: How many time zones does Russia have?

Question 5: In miles, how far does the average American drive each year?

Question 6: What percent of the world's population lives in the U.S.?

Question 7: On average, what percent of a watermelon's weight comes from water?

Wits and Wagers...

Question 8: What percentage of Americans say that reading is their favorite leisure-time activity?

Question 9: In feet and inches, how tall was the tallest human in recorded history?


- Enter this as feet.inches e.g., 3.0, 5.11

Question 10: In what year was an ATM machine first installed in the U.S.?


Please enter your interval on the Canvas quiz/survey

Enter your range of estimates for each question below as two numbers separated by a comma

- E.g., 10.2, 50.7



Answers...



Wits and Wagers...

Question 1: What percent of the world's surface is water?

- 71%

Question 2: How many floors does the leaning tower of Pisa have?

- 8

Question 3: What year was the parking meter invented?

- 1935

Wits and Wagers...

Question 4: How many time zones does Russia have?

- 11

Question 5: In miles, how far does the average American drive each year?

- 13,476

Question 6: What percent of the world's population lives in the U.S.?

- 4.27%

Question 7: On average, what percent of a watermelon's weight comes from water?

- 92%

Wits and Wagers...

Question 8: What percentage of Americans say that reading is their favorite leisure-time activity?

- 35%


Question 9: In feet and inches, how tall was the tallest human in recorded history?

- 8' 11"

Question 10: In what year was an ATM machine first installed in the U.S.?

- 1969

100% confidence intervals



There is a tradeoff between the **confidence level** (percent of times we capture the parameter) and the **confidence interval size**

Note

For any given confidence interval we compute, we don't know whether it has really captured the parameter

But we do know that if we do this 100 times, 95 of these intervals will have the parameter in it


(for a 95% confidence interval)

Confidence Intervals

For a **confidence level** of 90%...

90% of the **confidence intervals** will have the parameter in them


Right???



Computing confidence intervals

We will learn how to compute confidence intervals over the next few classes, but first let's review some material...

Sampling distributions are usually normal



So 95% of the sample statistics will fall within ± 2 Standard Errors from the mean

Thus, if we knew the SE, we could calculate a 95% confidence interval!

- Why is this true?

The problem

Unfortunately, we don't know the Standard Error ☹️

Q: Could we repeat the sampling process many times to estimate it?

A: No, we can't do an experiment that many times

We're just going to have to pick ourselves up from the bootstraps!

Estimate SE with \hat{SE}

Then use $\bar{x} \pm 2 \cdot \hat{SE}$ to get the 95% CI

Next class

Using the bootstrap to estimate the SE...

